

Abhinandan

+91-8130635690 | abhinandan@abhinandan.one

abhinandan.one | linkedin.com/in/abhibuilds | github.com/awesome-pro

EXPERIENCE

Founding Software Engineer

Sept 2025 – Present

Browser

Remote

- Built Browser's Chrome MV3 recorder + CDP-native browser automation agent, achieving 95%+ precise AX/DOM element capture with cross-iframe support, obstruction checks, and real mouse/key/upload execution.
- Built a smart streaming ReAct loop across FastAPI + extension with SSE tool execution, multi-tab orchestration, safe parallelism, abort/continue, and audit logs.
- Reduced automation LLM spend by roughly 67% using compact recording traces, context-window compression, prompt caching, and model-routing across GPT-5, Claude Sonnet & Haiku.
- Shipped a zero-LLM replay engine: recordings run as variable-driven tool-call templates, with a stateful AI fallback that resumes mid-run on failure.
- Shipped self-healing docs that auto-repair on UI drift. Haiku → Sonnet diff triage, LLM-free replay of intact steps, and a CDP agent that fixes only what changed.

Founding Software Engineer

Jan 2025 – Jul 2025

Cynos Nexus

Noida, India

- Built the AI knowledge base service using FastAPI, LangChain, and vector retrieval pipelines, powering customer-facing search workflows.
- Developed document-ingestion pipelines using Google Cloud Vision, XLSX processing, and BullMQ workers, enabling automated extraction of customer data from spreadsheets and scanned records.
- Implemented backend infrastructure for retrieval-augmented AI workflows, including document ingestion, structured data processing, asynchronous workers, Redis-backed queues, and production API serving.
- Automated containerized deployment and CI/CD workflows for AI services using Docker, GitHub Actions, Nginx reverse proxying, and cloud infrastructure.

Contract Software Engineer

Sept 2024 – Dec 2024

Etkin.ai

Remote, Turkiye

- Built a LangChain + pgvector knowledge base powering AI-assisted document search and retrieval workflows, improving query accuracy by 15%.
- Developed scalable data-ingestion pipelines using bulk CSV processing and Celery workers, reducing processing time by 40%.
- Engineered a production PDF generation system transforming structured AI outputs and dynamic JSON reports into enterprise-grade documents.
- Automated deployment of AI services using Docker, GitHub Actions, and AWS EC2, establishing reliable CI/CD workflows for production environments.

Open Source Contributor

June 2024 – Aug 2024

HeroUI

Y Combinator S24 startup

- Received a **personal offer from the CEO** to join HeroUI (prev. NextUI) after making **open-source contributions**.
- Resolved 10+ bugs & delivered 7+ feature enhancements in core components including Calendar, Table and Pagination.

SWE Intern

Jan 2024 – Apr 2024

SkilledUp

Noida

OPEN SOURCE PROJECTS

AgentFlow-Pro (*Research*) | *PyTorch, TRL, DAPO, PRM, PEFT/LoRA, FastMCP* | github.com/awesome-pro/agentflow-pro

- Rebuilt and extended the ICLR 2026 *AgentFlow* architecture as a local Qwen3-8B Planner → Executor → Verifier → Memory loop with grammar-constrained planning, Tavily search, and sandboxed Python/SymPy tools.
- Trained the planner with DAPO + a learned Qwen3-0.6B Process Reward Model, replacing the paper's GRPO setup with a custom dynamic-sampling stage and full A40 training pipeline.
- Built a leakage-free, quantization-matched evaluation flow and measured +5.0 pts on GPQA-Diamond, improving from 40.0% to 45.0%.

smartmemo | *PyTorch, sentence-transformers, FAISS, SQLite* | github.com/awesome-pro/smartmemo

- Built a semantic memory/cache for LLM agents where embedding retrieval proposes candidates and a learned pair-equivalence classifier decides whether memory reuse is safe.
- Trained classifier-v2 on 16,576 pairs across 9 domains, improving precision by +30 points at equal recall versus cosine similarity, with feedback export, bad-hit detection, gated retraining, and WAL-backed SQLite hardening.

agenteval | *AsyncIO, OpenAI/Anthropic/LangChain adapters* | github.com/awesome-pro/agenteval

- Agent evaluation toolkit replacing exact-match asserts with repeated-run pass-rate tests; traces tool calls/timing/steps, supports collect-then-raise behavioral assertions, OpenAI/Anthropic/LangChain adapters, and Typer CLI JSON reports for CI gates.

guardloop | *OpenAI/Anthropic SDKs, OpenTelemetry, LangGraph* | github.com/awesome-pro/guardloop

- Production guardrail runtime for async agents: pre-flight Decimal cost/token/time/tool-call budgets, per-tool circuit breakers, verifier feedback retry loop, OpenTelemetry GenAI spans, and no-rewrite LangGraph/OpenAI Agents SDK adapters.

TECHNICAL SKILLS

ML Engineering: PyTorch, Hugging Face, sentence-transformers, embeddings, classifier training, RAG, FAISS, Pinecone, Qdrant, ChromaDB, pgvector

Agents & Frameworks: LangGraph, OpenAI Agents SDK, LangChain, LlamaIndex, CrewAI

LLM Engineering: Fine-tuning (LoRA, QLoRA, PEFT), RLHF / DPO / GRPO / DAPO / PRM, quantization (AWQ, GPTQ, GGUF), inference engines (vLLM, SGLang, Ollama, TGI), context engineering, KV-cache strategies, streaming

Evals & Reliability: OpenTelemetry, Langfuse, Arize Phoenix, LangSmith, statistical pass-rate evaluation, AIME / GPQA benchmarks

Backend & Infrastructure: FastAPI, Node.js, GraphQL, WebSockets, AWS, GCP, Redis, Docker

Languages: Python, TypeScript, C++, SQL, Bash

AWARDS & RECOGNITION

International Youth Math Challenge Gold Honour | **Top 1% TypeScript Engineer Globally** (Algora) | **Amazon ML Summer School 2025** | 2nd Place, Outlier AI Hackathon

Reliance Foundation Scholar | **HDFC Badhate Kadam Scholar**

EDUCATION

Dr. APJ Abdul Kalam Technical University

B.Tech, Computer Science | Nov 2022 – Apr 2026